

# Underfitting




# Underfitting

## What is Underfitting ?

Underfitting corresponds to a model that :

- Is too simple and inflexible to learn
- Cannot model (learn) the training data
- Cannot generalize to treat new data




A model that has a poor performance.  
This bad performance is detected through the evaluation metrics.

# Underfitting

How to reduce/avoid Underfitting ?


Any ideas ?



# Underfitting

## How to reduce/avoid Underfitting ?

- Increase model complexity
- Increase data features (independent variables)
- Improve data (remove noise, remove redundancy)
- Add more data
- Increase the number of epochs
- Increase training duration





# Overfitting



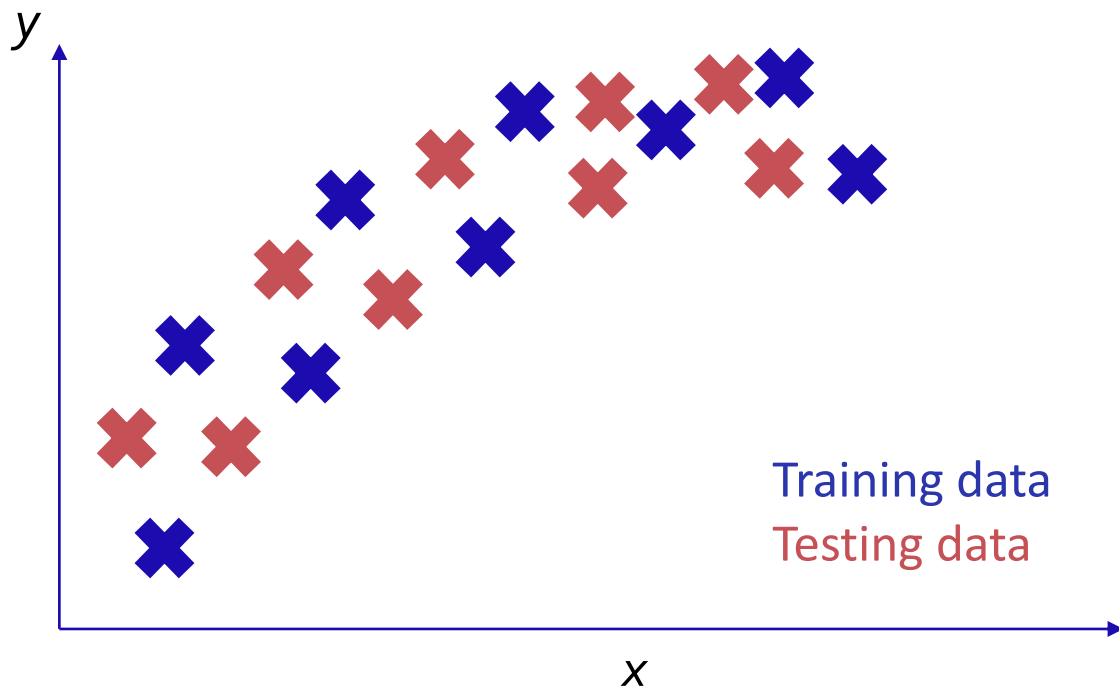
# Overfitting

Bias-Variance Tradeoff :




# Overfitting

Bias-Variance Tradeoff :



Trained Model 2

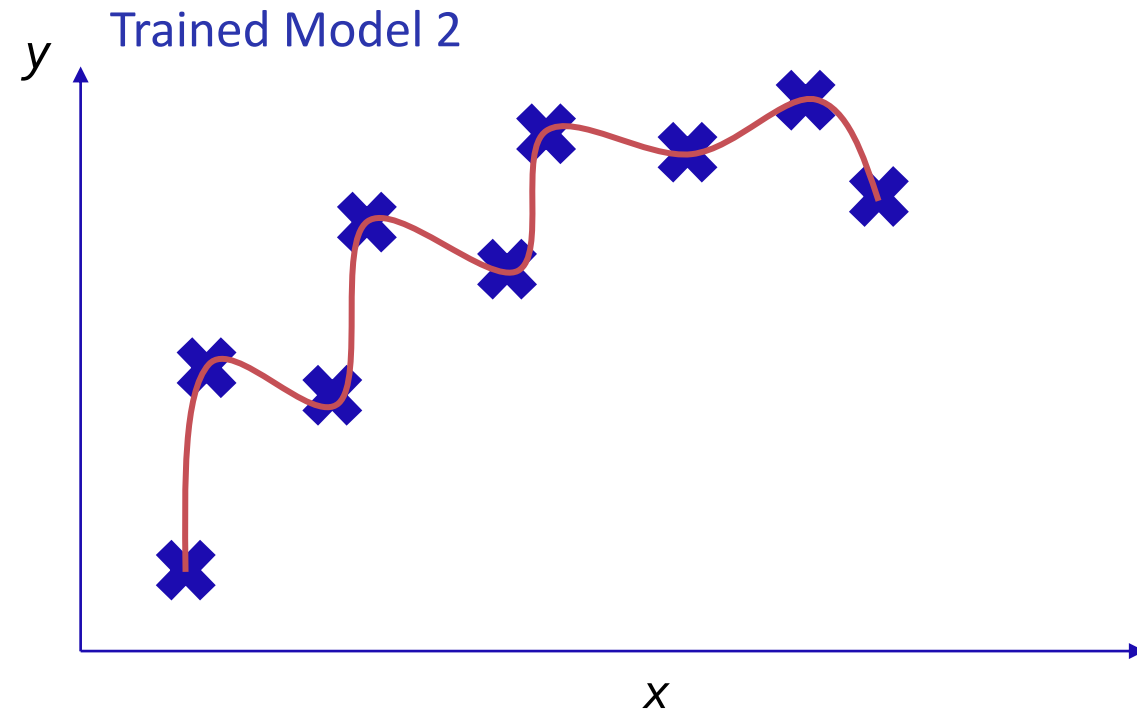
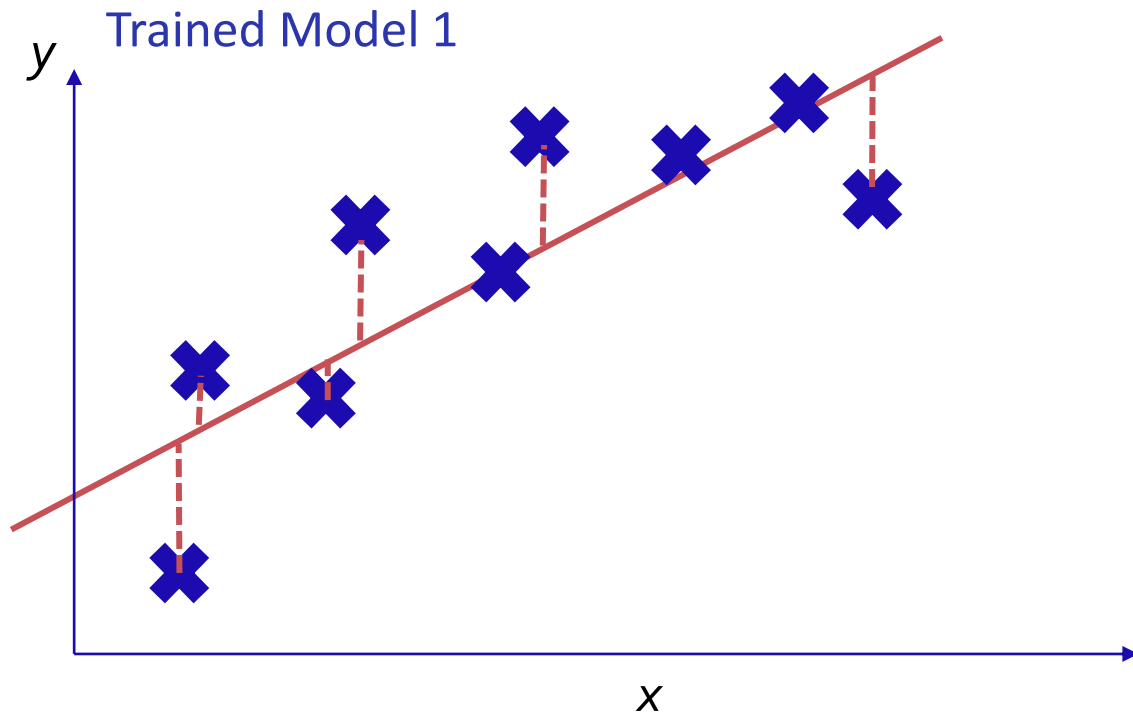


# Overfitting

Bias-Variance Tradeoff :

We compare both trained models using RMSE (Root-Mean-Square Error) :

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2}$$







# Overfitting

Bias-Variance Tradeoff :

We compare both trained models using RMSE :




# Overfitting

Bias-Variance Tradeoff :


We compare both trained models using the test data:

$RMSE\ 1 < RMSE\ 2 \rightarrow$  Model 1 has a better performance with new data !

Trained Model 1





Trained Model 2




# Overfitting


The incapability of a model to reproduce data exactly is called **bias**



The different 'fits' of a trained model to different datasets is called the **variance**




# Overfitting





Relatively high bias  
Relatively low variance

→ Predictable model



Very low bias  
Very high variance



→ Unpredictable model



# Overfitting

Overfitting :

The ideal model has both the bias and the variance low




Since this model has a perfect fit to the training data and a bad performance with new data

→ The model is overfit

# Overfitting

How to reduce/avoid Overfitting ?


Any ideas ?



# Overfitting

## How to reduce/avoid Overfitting ?

- Cross-validation to tune hyperparameters
- Use more data for training
- Feature selection (removing irrelevant features)
- Early stopping
- Regularisation

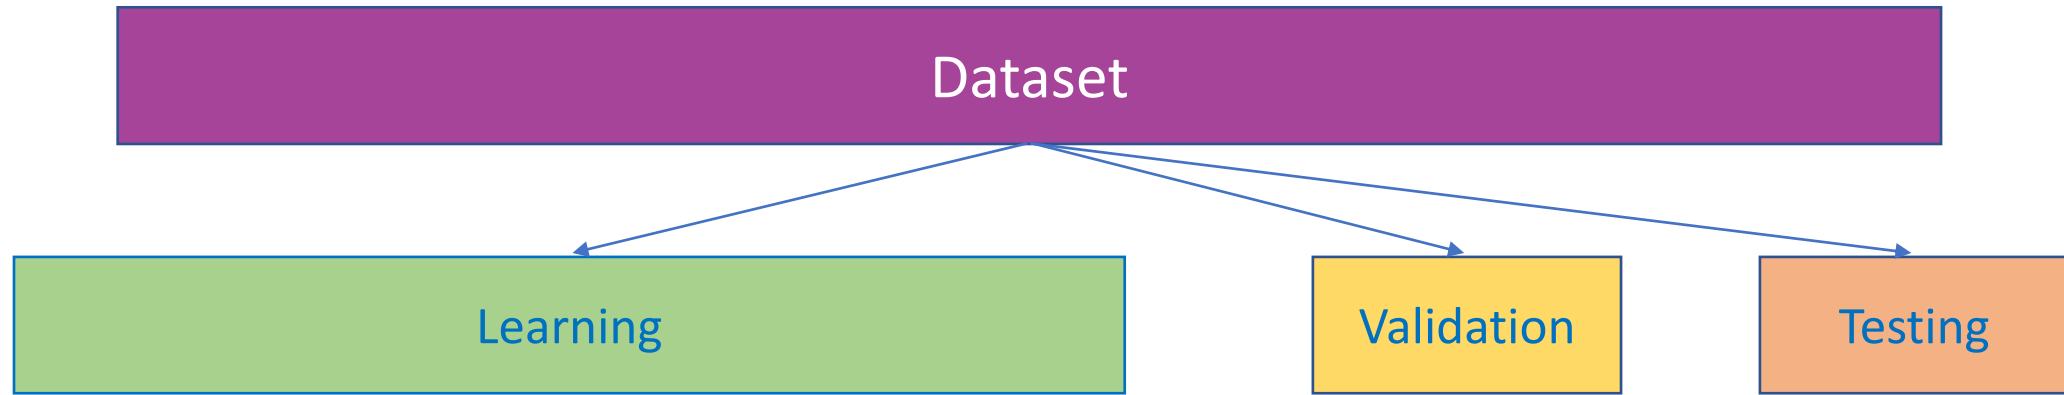


# Learning, Validation & Testing





# Learning, Validation & Testing



Due to overfitting risk, the learning loss is not a good indication of a well-trained model

==> We need validation and training datasets

Note : These three datasets are separate and contain no repeated data

# Learning, Validation & Testing

**Learning (Training) dataset** : The data used to train the model (adjust the weights)

- During each epoch, our model will be trained over the entire learning dataset
- Labeled (loss & accuracy in each epoch can be provided)

**Validation dataset** : The data used to validate our model during training

- After weights are adjusted, the model is validated using this dataset
- Does not contribute in the weight adjustment (learning)
- Helps us choose the best hyperparametres for our model
- Helps us detect overfitting
- Labeled (loss & accuracy in each epoch can be provided)

# Learning, Validation & Testing

**Test dataset** : The data used to test the model after it has been trained and validated

- Should not be labeled (the same condition as in using our deployed model in real life)

The ultimate goal of machine learning and deep learning is to build models that are able to generalize and perform well on new unseen data.

# Learning, Validation & Testing

## Hyperparameters :

- Variables set before the model's training
- Help in model selection
- Include parameters that :
  - specify the network structure (number of hidden layers)
  - determine how the network is trained (number of epochs, batch size, learning rate)

# Learning, Validation & Testing

## Définitions :



### **Batch size :**

The size of the training set

### **Hidden layers :**

Layers between the input and the output layers in the neural network

### **Epoch :**


A pass over the entire training database

### **Learning rate :**

A parameter that controls the rate at which the neural network learns (ranges between 0.0 and 1.0)

# Learning, Validation & Testing

Cross-Validation :



**Cross-validation** trains the model on subsets of the available data and evaluates it on the complementary subset that was not used for training.

# Learning, Validation & Testing

Cross-Validation :

Training

Validation



**K-fold cross-validation :**

- The data is split into  $k$  subsets/folds.
- Training is done on  $k-1$  subsets
- Evaluation is done on the left-out subset.
- This process is repeated  $k$  times with a different subset left for evaluation each time.

# Learning, Validation & Testing

## Cross-Validation :

- Helps detect overfitting
- Is used to optimize hyperparameters of a model
- Is used for model selection





# Learning, Validation & Testing

Example : Model selection with k-fold cross-validation (with  $k = 4$ )

Model 1

% Quality of performance
85%
87%
90%
88%

Model 2

% Quality of performance
87%
88%
91%
89%

Model 3

% Quality of performance
87%
88%
91%
89%

Model 4

% Quality of performance
87%
88%
91%
89%




Overall performance  
= 87,5%  
(average of performances)

88,75%

93,5%

97,0%

# Regularisation



# Regularisation

Regularisation :

- is a technique that introduces slight modifications to the learning algorithm for it to generalize better
- also improves the model's performance on new unseen data
- helps reduce overfitting by reducing the model's complexity
- introduces a small bias to our model that, in turn, reduces the variance significantly

# Regularisation

## 1. L1 Regularisation (Lasso Regression) :

$$L1Regularisation = Lossfunction + \lambda \sum_{i=1}^n |w_i|$$

Regularisation term that adds bias to the model to avoid perfect fitting

# Regularisation

## 1. L1 Regularisation (Lasso Regression) :

$$L1Regularisation = Lossfunction + \lambda \sum_{i=1}^n |w_i|$$

Regularisation term that adds bias to the model to avoid perfect fitting

$\lambda$  determines how strongly the regularization influences the network's training.

- If  $\lambda = 0$  --> no regularization at all
- If  $\lambda = 1,000,000,000$  --> very high regularization (too much bias, model hardly learns anything)

Reasonable choices of the regularization strength  $\lambda$  (0.001, 0.01, 0.1, 1.0, etc.)

The best  $\lambda$  can be chosen using cross validation

# Regularisation

## 2. L2 Regularisation (Ridge Regression) :

$$L2Regularisation = Lossfunction + \lambda \sum_{i=1}^n w_i^2$$

Regularisation term that adds bias to the model to avoid perfect fitting

The best  $\lambda$  can be chosen using cross validation

# Regularisation

## L1 versus L2 Regularisation :

L1 Regularisation shrinks less important parameters to 0

L2 Regularisation shrinks less important parameters close to 0 but not equal to 0

L2 Regularisation penalises larger weights more severely :

\*\* Example : A weight of 10 get a penalty = 10 in L1 Regularisation and a penalty = 100 in L2 Regularisation

# Regularisation

## 3. Dropout Regularisation :

Dropout regularisation :

- is a computationally cheap type of regularisation
- Produces very good results
- Is the most frequently used regularisation technique in deep learning




# Regularisation

## 3. Dropout Regularisation :

Dropout regularisation : At every iteration, the algorithm randomly selects some nodes and removes all connections (inputs & outputs) to them.


Each iteration has a different set of dropped out nodes



# Regularisation

## 3. Dropout Regularisation :

The probability of choosing the number of nodes to drop out is a hyper parameter that we can assign to the dropout function.



Dropout can be applied to both the hidden and the input layers

Dropout is usually preferred for large neural networks

# Regularisation

## 4. Early Stopping :

As soon as the performance on the validation dataset starts to decrease (loss starts to increase), the training is stopped immediately, hence the **early stopping**.




# Regularisation

## 4. Early Stopping :

Example :

The training is stopped at the dotted line



# Regularisation

## 5. Data Augmentation :

Data augmentation corresponds to increasing the size of the training dataset

\*\* Example : We can introduce variations in images : rotation, flipping, rescaling, shifting

# Regularisation

## 5. Data Augmentation :

### Base Augmentations

Geometry based



rotate



shear



vertical-flip



horizontal-flip



crop



crop-and-pad



Perspective-transform



Elastic-transformation

Color based



sharpen



brighten




Gamma-contrast



invert

Noise / occlusion



gaussian-blur



additive-gaussian-noise




translate-x



translate-y



coarse-salt



super-pixel



emboss

Weather



clouds



fog



snow-flakes



Fast-snowy-landscape

# Artificial Neural Networks

by

*Dr. Vera DAMERJAN PIETERS*

# Thank you

